

Hackathon KAUST GPU 2020

Marcin Rogowski, Suha Kayum & Vincent Etienne

KAUST & Saudi Aramco

November 2020

- 1 hpcscan
 - Overview
 - Compilation and validation
- 2 Test platforms
 - Shaheen II (KAUST)
- 3 Test Case Grid
- 4 Test Case Comm
- 5 Test Case FD_D2
- 6 Test Case PropaAc2
- 7 Conclusions and next steps
- 8 Acknowledgements

- 1 **hpcscan**
 - Overview
 - Compilation and validation
- 2 Test platforms
 - Shaheen II (KAUST)
- 3 Test Case Grid
- 4 Test Case Comm
- 5 Test Case FD_D2
- 6 Test Case PropaAc2
- 7 Conclusions and next steps
- 8 Acknowledgements

hpcscan is a C++ code for benchmarking HPC kernels (mainly for solving PDEs with FDM)

- Simple code structure based on individual test cases
- Easy to add new test cases
- Main class is Grid: multi-dimension (1, 2 & 3D) Cartesian grid
- Hybrid MPI/OpenMP parallelism
- All configuration parameters on command line
- Support single and double precision computation
- Compilation with standard Makefile
- No external libraries
- Follows C++ Google style code

hpcscan embeds several test cases

Current version 1.0

- General operations on grids
- Memory operations
- MPI communication
- FD computation
- Basic wave propagator

Possible additions for future versions

- Operations on matrices full and sparse
- FFT
- IO
- Compression

Compilation and validation

Compiling hpcscan

go to `./build` and `make` (by default compilation with single precision float)

To compile with double precision float, make `precision=double`

Validating hpcscan

go to `./script` and `sh runValidationTests.sh`

Table: `runValidationTests.sh`¹

Machine	Compiler	Single prec.	Double prec.
Mars	g++ 9.3.0	764 PASS / 0 FAIL / 0 ERR / 20 WARN	764 PASS / 0 FAIL / 0 ERR / 20 WARN
Shaheen	icpc 19.0.5.281	764 PASS / 0 FAIL / 0 ERR / 20 WARN	764 PASS / 0 FAIL / 0 ERR / 20 WARN

Numbers can differ due to availability of features depending on the platforms

¹Updated Nov 25, 2020

- 1 hpcscan
 - Overview
 - Compilation and validation
- 2 **Test platforms**
 - **Shaheen II (KAUST)**
- 3 Test Case Grid
- 4 Test Case Comm
- 5 Test Case FD_D2
- 6 Test Case PropaAc2
- 7 Conclusions and next steps
- 8 Acknowledgements

Test platform - Shaheen II (KAUST)

Machine Shaheen II / Cray XC40

- Computing nodes Intel Haswell 2.3 Ghz dual socket (16 cores / socket)
- RAM 128 GB with Peak memory BW 136.5 GB/s
- Peak performance Single Prec. 2.36 TFLOP/s / Double Prec. 1.18 TFLOP/s
- Interconnect Cray Aries with Dragonfly topology
 - 60 GB/s optical links between groups
 - 8.5 GB/s copper links between chassis
 - 3.5 GB/s backplane within a chassis
 - 5 GB/s PCIe from node to Aries router



- 1 hpcscan
 - Overview
 - Compilation and validation
- 2 Test platforms
 - Shaheen II (KAUST)
- 3 **Test Case Grid**
- 4 Test Case Comm
- 5 Test Case FD_D2
- 6 Test Case PropaAc2
- 7 Conclusions and next steps
- 8 Acknowledgements

Test Case Grid - Description

- Fill grid ($W = \text{coef}$)
- Max. err. grid W
- L1 err. grid W
- Get min. grid W
- Get max. grid W
- Update pressure (used in propagator)
- Small Grid size 500 MB ($500 \times 500 \times 500$ points)
- Medium Grid size 4 GB ($1000 \times 1000 \times 1000$ points)

Test Case Grid - Results

- Machine: shaheen
- 1 node / 32 threads
- Baseline kernel

Table: Bandwidth GB/s ²

Grid	Fill	Max. err.	L1 err.	Get max.	Get min.	Update Pres.
Small	58	72	122	125	125	119
Medium	54	91	124	127	127	120

Table: Bandwidth GPoints/s

Grid	Fill	Max. err.	L1 err.	Get max.	Get min.	Update Pres.
Small	14.4	9.0	15.2	31.3	31.2	6.0
Medium	13.4	11.4	15.5	31.8	31.8	6.0

Reproduce results with `./script/testCase_Grid/runSmallGridShaheen.sh` and
`./script/testCase_Grid/runMediumGridShaheen.sh`
Elapsed time 5 and 7 sec.

²Updated Nov 26, 2020

Machine: Shaheen

- L1 Err., Get Min & Max: 125 GB/s close to peak BW (92 % Peak Mem. BW)
- Low perf for Fill: 54-58 GB/s (40-43 % Peak Mem. BW)
- Max Err. 72-91 GB/s (53-67 % Peak Mem. BW)
- Pressure update 6 GPoint/s (120 GB/s, 88 % Peak Mem. BW)

- 1 hpcscan
 - Overview
 - Compilation and validation
- 2 Test platforms
 - Shaheen II (KAUST)
- 3 Test Case Grid
- 4 Test Case Comm**
- 5 Test Case FD_D2
- 6 Test Case PropaAc2
- 7 Conclusions and next steps
- 8 Acknowledgements

Measure MPI communication bandwidth

MPI point to point communication

- Send with MPI_Send from proc X to proc 0 (Half-duplex BW)
- Send and receive with MPI_Sendrecv between proc X and proc 0 (Full-duplex BW)

MPI collective communication

- Exchange of halos used in FD kernel with MPI_Sendrecv
- Grid size $1000 \times 1000 \times 1000$
- Domain decomposition with $N1 \times N2 \times N3$ subdomains

Test Case Comm - Results

- **Machine: Shaheen**
- 8 MPI processes (1 per computing node)
- Baseline kernel

Table: Bandwidth GB/s ³

MPI#1	MPI#2	Send	Sendrecv	Halo exch.	Comm. size	Subdomains
0	1	8.5	15.3	-	47 MB	-
0	2	8.3	15.3	-	47 MB	-
0	3	8.6	15.3	-	47 MB	-
0	4	8.5	15.3	-	47 MB	-
0	5	8.2	15.3	-	47 MB	-
0	6	8.5	15.3	-	47 MB	-
0	7	8.6	15.3	-	47 MB	-
All	All	-	-	5.0	128 MB	1 4 2
All	All	-	-	5.1	128 MB	1 2 4
All	All	-	-	2.0	96 MB	2 2 2

Reproduce results with `./script/testCase_Comm/runTestShaheen.sh`

Elapsed time 9 seconds

³Updated Sep 19, 2020

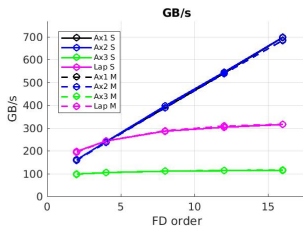
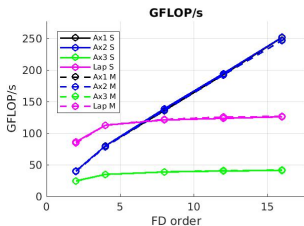
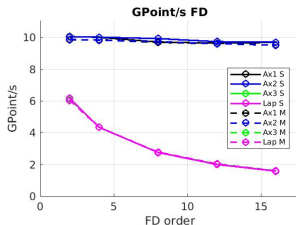
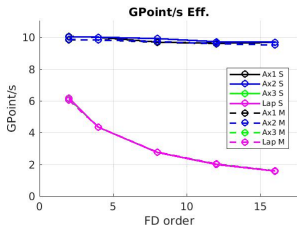
- 1 hpcscan
 - Overview
 - Compilation and validation
- 2 Test platforms
 - Shaheen II (KAUST)
- 3 Test Case Grid
- 4 Test Case Comm
- 5 Test Case FD_D2**
- 6 Test Case PropaAc2
- 7 Conclusions and next steps
- 8 Acknowledgements

Test Case FD_D2 - Description

- Computation of second order derivatives with finite-difference stencil
- Directional derivatives
 - Axis 1 $W = \partial_{x1}^2(U)$
 - Axis 2 $W = \partial_{x2}^2(U)$
 - Axis 3 $W = \partial_{x3}^2(U)$
- Laplacian
 - For 2D grids $W = \Delta(U) = \partial_{x1}^2(U) + \partial_{x2}^2(U)$
 - For 3D grids $W = \Delta(U) = \partial_{x1}^2(U) + \partial_{x2}^2(U) + \partial_{x3}^2(U)$
- Stencil order 2, 4, 8, 12 & 16
- Grid size
 - Small $500 \times 500 \times 500$
 - Medium $1000 \times 1000 \times 1000$

Test Case FD_D2 - Results

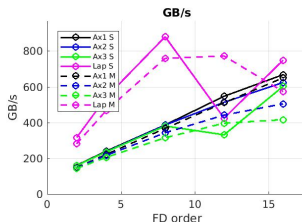
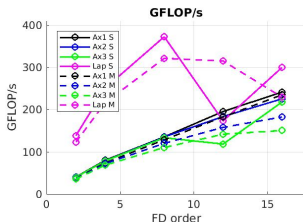
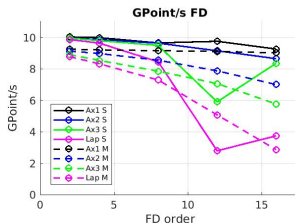
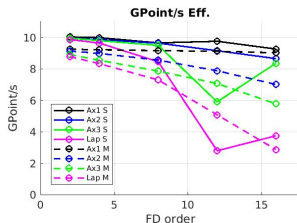
- machine Shaheen / 1 node with 32 threads / Baseline kernel ⁴
- ./script/testCase_FD_D2/runSmallGridShaheen.sh & runMediumGridShaheen.sh



⁴ Updated Sep 26, 2020

Test Case FD_D2 - Results

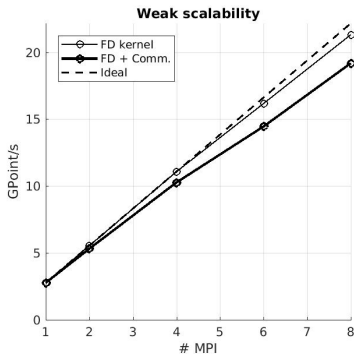
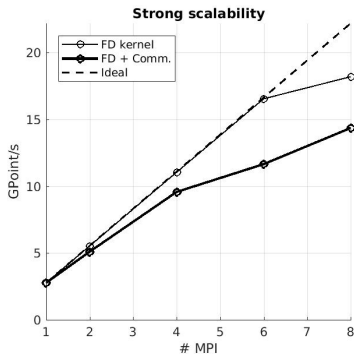
- machine Shaheen / 1 node with 32 threads / Cache blocking kernel ⁵
- ./script/testCase_FD_D2/runSmallGridShaheen.sh & runMediumGridShaheen.sh



⁵ Updated Sep 26, 2020

Test Case FD_D2 - Results

- machine Shaheen
- 1 to 8 nodes with 32 threads/node
- Baseline kernel ⁶
- Strong scalability: Grid $1000 \times 1000 \times 1000$ (4 GB)
- Weak scalability: Grids from 4 GB (1 proc) to 32 GB (8 proc)
- 3D Laplacian O8



⁶ Updated Sep 26, 2020

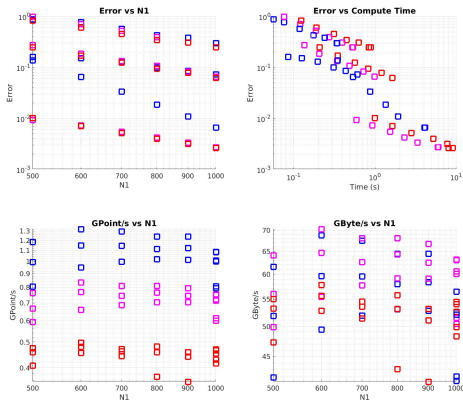
machine Shaheen

- Large benefit of cache blocking
- Significant effect of grid dimension and index (very bad performance for n3 without cache blocking)
- Min BW 50 GFLOP/s (∂_{x3}^2 O2) = 2 % peak BW [apparent Mem. BW 150 GB/s]
- Max BW 370 GFLOP/s (Δ O8) = 16 % peak BW [apparent Mem. BW 900 GB/s]
- Apparent Mem. BW 150-900 GB/s (110-660 % Peak Mem. BW) = shows data in-cache effect
- Typical stencils of interest for geophysical applications
 - Δ O4 BW = 8-10 GPoint/s
 - Δ O8 BW = 7-9 GPoint/s
 - Δ O12 BW = 3-5 GPoint/s
- Parallel efficiency with 8 nodes 55 to 86 % (depends on workload on Shaheen)

- 1 hpcscan
 - Overview
 - Compilation and validation
- 2 Test platforms
 - Shaheen II (KAUST)
- 3 Test Case Grid
- 4 Test Case Comm
- 5 Test Case FD_D2
- 6 Test Case PropaAc2**
- 7 Conclusions and next steps
- 8 Acknowledgements

Test Case PropaAc2 - Results

- machine Mars / preliminary results ⁷
- Eigen mode - 1D model
- FD: Black O2, Blue O4, Pink O8, Red O12 / Square=Baseline
- `./paramAnalysis/propaAccuracy/runMars.sh`



⁷Updated Nov 5, 2020

- 1 hpcscan
 - Overview
 - Compilation and validation
- 2 Test platforms
 - Shaheen II (KAUST)
- 3 Test Case Grid
- 4 Test Case Comm
- 5 Test Case FD_D2
- 6 Test Case PropaAc2
- 7 Conclusions and next steps**
- 8 Acknowledgements

Conclusions and next steps

TO DO

- 1 hpcscan
 - Overview
 - Compilation and validation
- 2 Test platforms
 - Shaheen II (KAUST)
- 3 Test Case Grid
- 4 Test Case Comm
- 5 Test Case FD_D2
- 6 Test Case PropaAc2
- 7 Conclusions and next steps
- 8 Acknowledgements

Acknowledgements

- KAUST ECRC and KSL for access and support on Shaheen II & Ibex